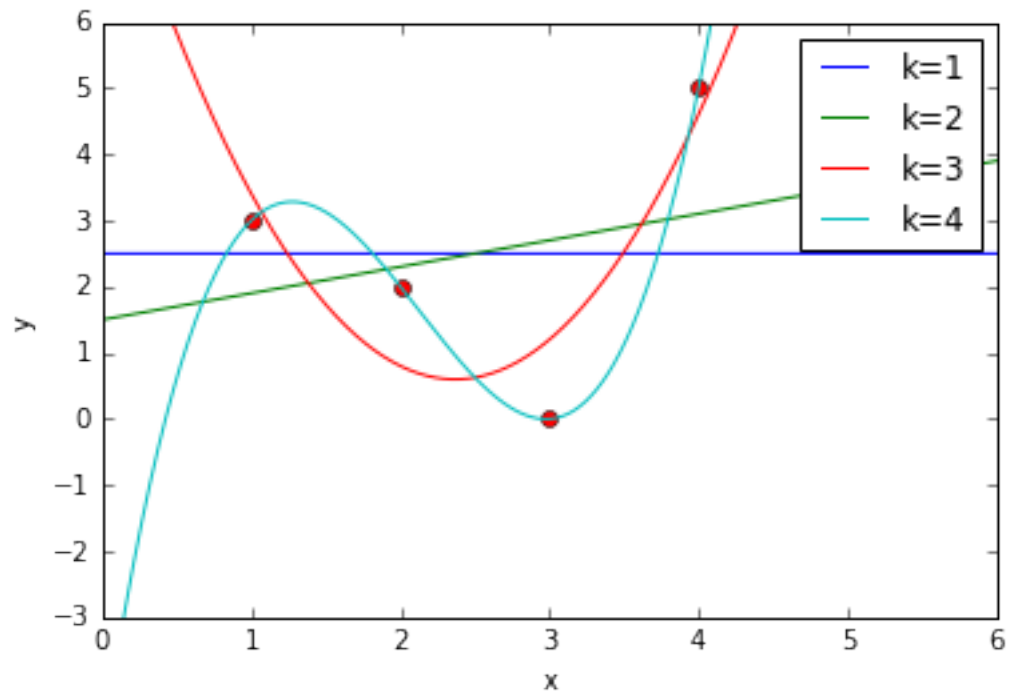


# Supervised Learning: Coursework 1

Matineh Akhlaghinia, Abdul Faiz Punakkath

November 14, 2018

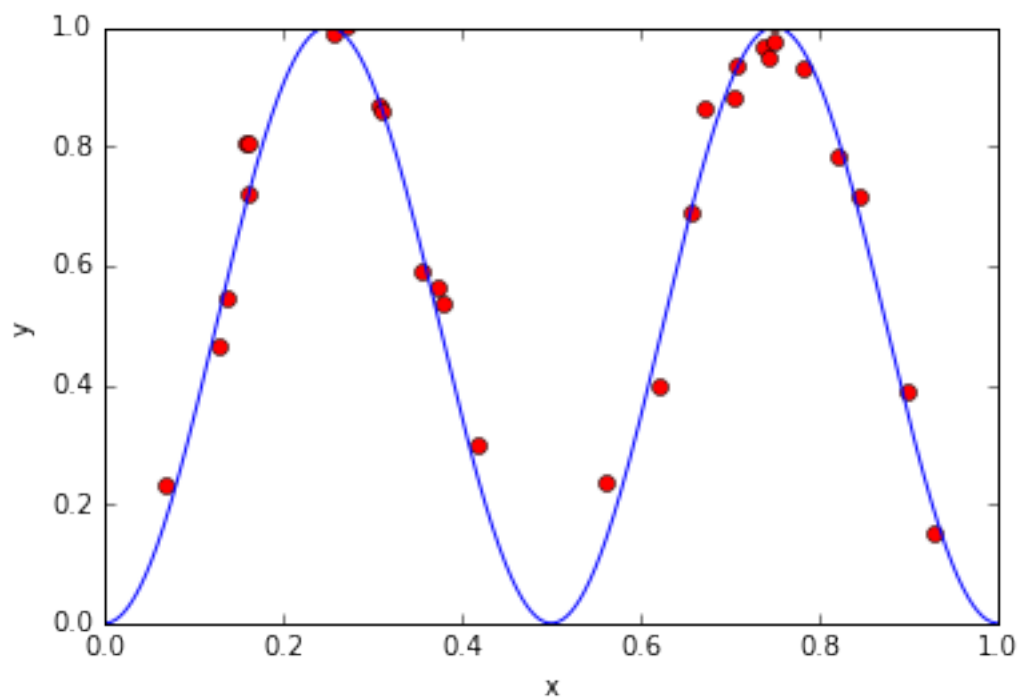
1. (a)



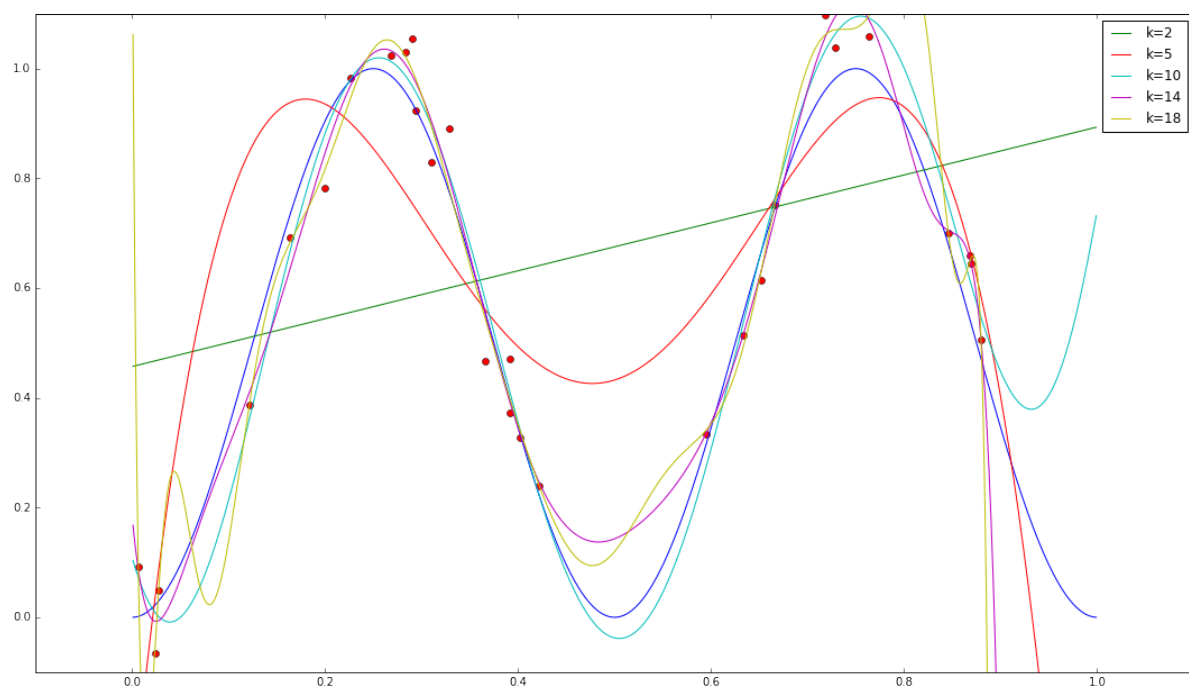
(b)  $k = 1 : 2.5$   
 $k = 2 : 1.5 + 0.4 * x$   
 $k = 3 : 9 - 7.1 * x + 1.5 * x^2$

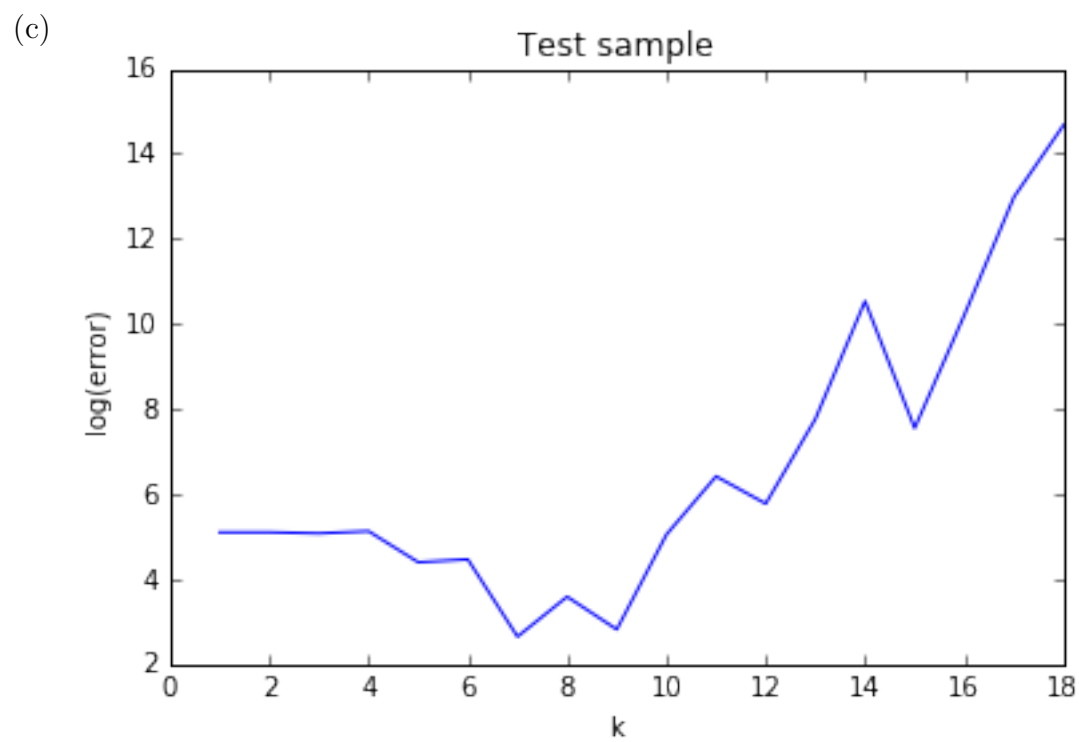
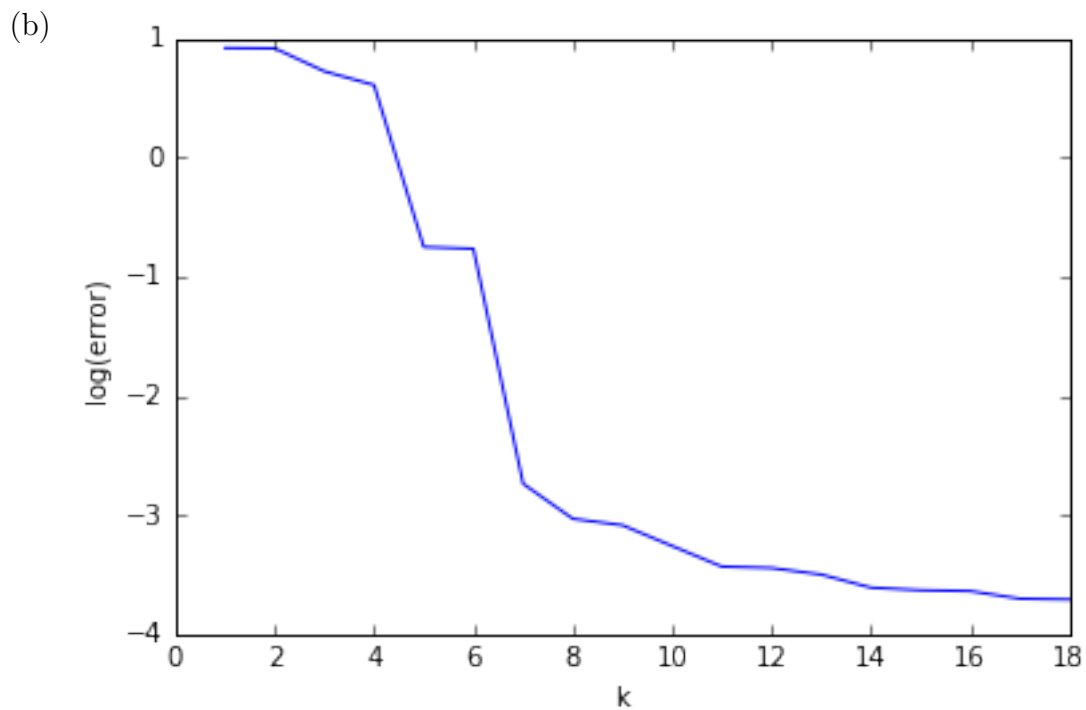
(c)  $k = 1 : 3.25$   
 $k = 2 : 3.0499999999999994$   
 $k = 3 : 0.7999999999999998$   
 $k = 4 : 6.589355141311112 * 10^{-27}$

2. (a) i.

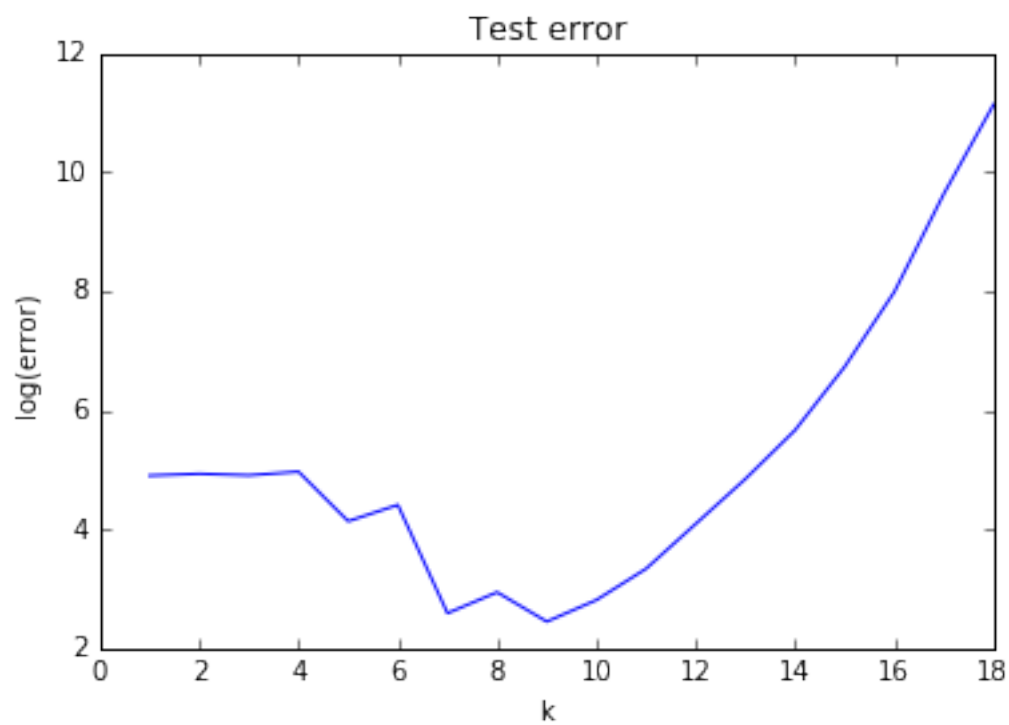
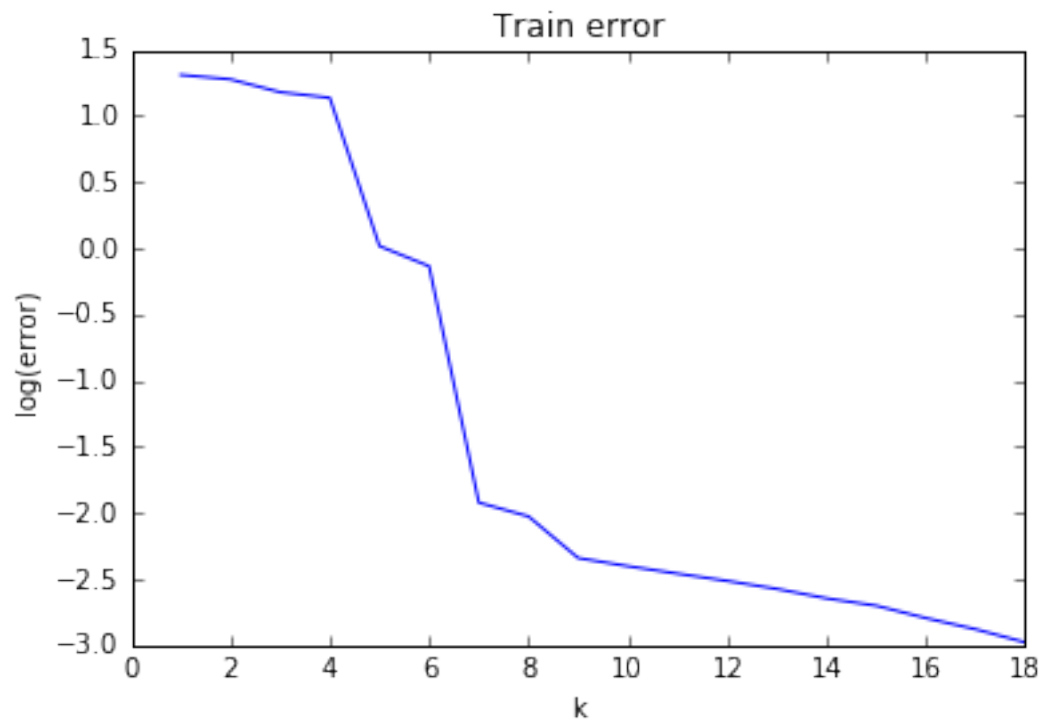


ii.

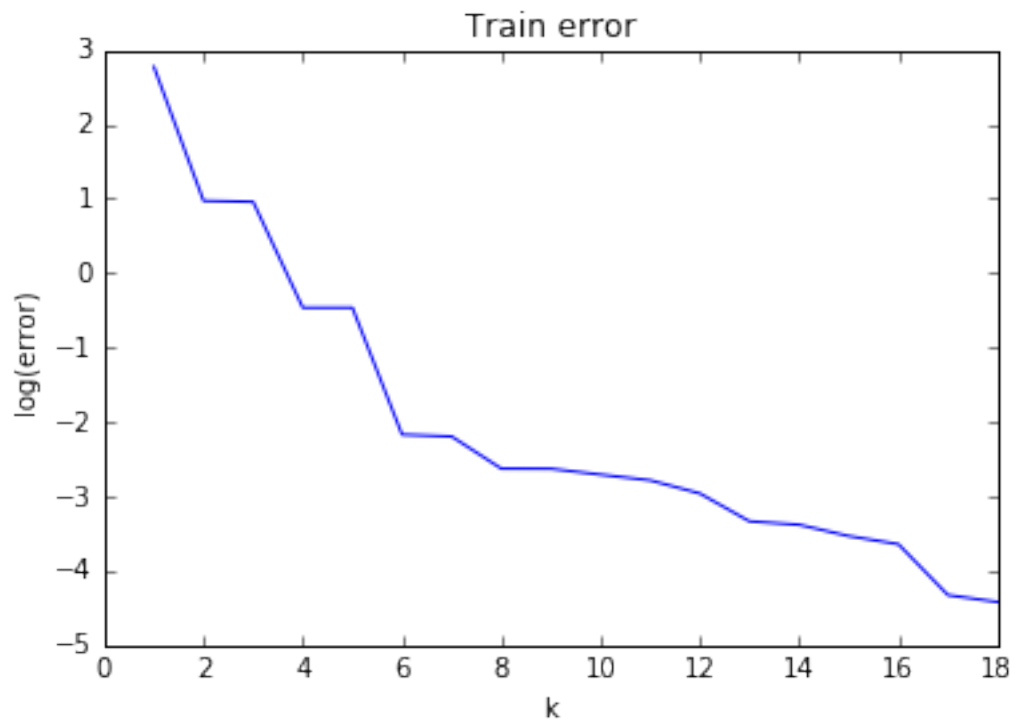


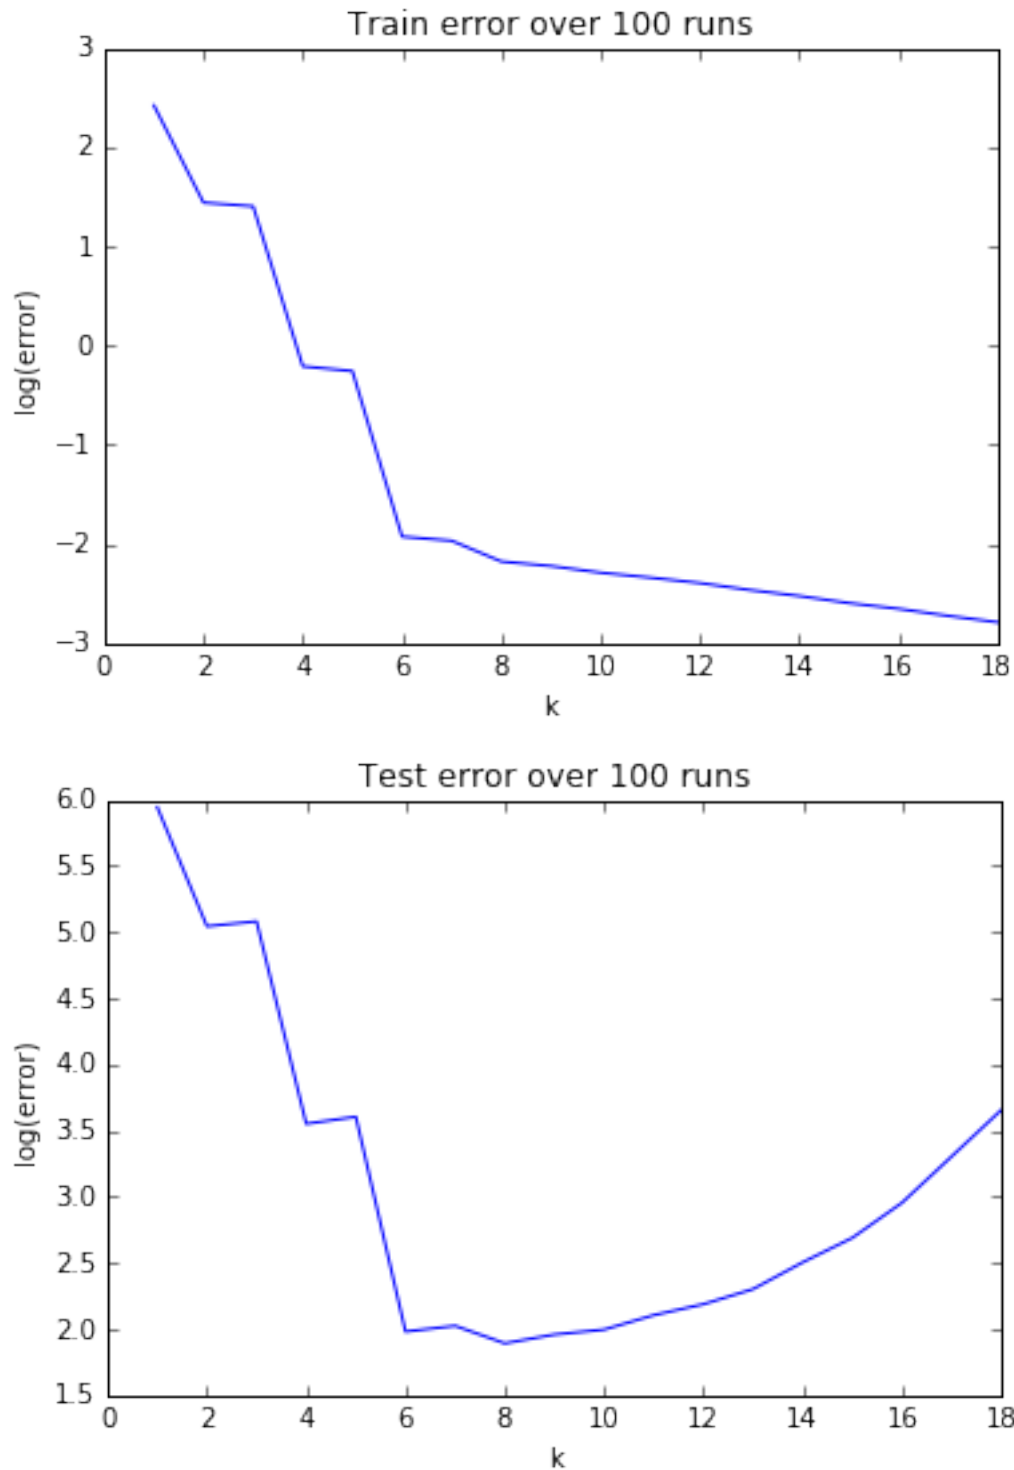


(d)



3.





4. (a)  $Train\ MSE = 74.94491097922847$   
 $Test\ MSE = 130.8907648716782$
- (b) Constant function in the above question will be the average of all the outputs. Since  $f(x) = (average\ of\ all\ the\ house\ prices)$  will achieve the minimum MSE for a constant function.

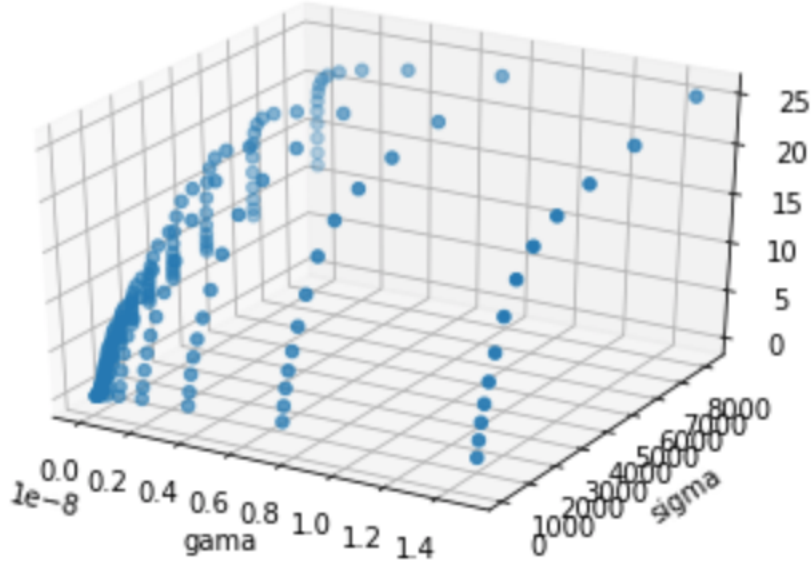
Attributes	Train MSE	Test MSE
1: CRIM	71.13968241182597	73.85021859717854
2: ZN	73.11604863223906	74.76428861359675
3: INDUS	64.47064841447379	65.39468687518753
4: CHAS	81.16463035498536	83.8633011280258
5: NOX	68.81917327442122	69.78392360937791
6: RM	43.78222803786089	43.740061995511425
7: AGE	73.88134430479434	70.01600700106209
8: DIS	78.85716721087809	80.20163857248215
9: RAD	72.34832512917372	72.08181955662278
10: TAX	65.5366146679282	67.04818858708646
11: PTRATIO	62.73194016874686	62.89616046358227
12: BLACK	73.68884487908194	78.03805991786986
13: LSTAT	38.44191281643897	38.89767608243805

(c)

(d)  $Train\ MSE = 21.11943099906709 // Test\ MSE = 25.155269280832382$ 

5. (a) The best sigma and gama are 128.0 and  $2^{-40}$ , respectively.

(b)

(c)  $Train\ MSE = 0.029$  $Test\ MSE = 2.23 \times 10^{-6}$

	<b>Method</b>	<b>MSETrain</b>	<b>MSETest</b>
	Naive Regression	84.88 $\pm$ 4.7	82.90 $\pm$ 9.4
	Linear Regression (attribute 1 (CRIM))	66.02 $\pm$ 4.1	64.12 $\pm$ 7.1
	Linear Regression (attribute 2 (ZN))	69.72 $\pm$ 4.8	68.47 $\pm$ 9.1
	Linear Regression (attribute 3 (INDUS))	59.83 $\pm$ 3.4	61.52 $\pm$ 6.8
	Linear Regression (attribute 4 (CHAS))	80.37 $\pm$ 3.2	83.49 $\pm$ 6.2
	Linear Regression (attribute 5 (NOX))	66.94 $\pm$ 4.3	71.82 $\pm$ 9.3
	Linear Regression (attribute 6 (RM))	36.75 $\pm$ 2.8	39.19 $\pm$ 5.3
(d)	Linear Regression (attribute 7 (AGE))	69.32 $\pm$ 5.7	71.55 $\pm$ 11.7
	Linear Regression (attribute 8 (DIS))	73.31 $\pm$ 6.4	78.22 $\pm$ 12.8
	Linear Regression (attribute 9 (RAD))	69.27 $\pm$ 4.8	67.87 $\pm$ 9.8
	Linear Regression (attribute 10 (TAX))	57.52 $\pm$ 2.6	58.17 $\pm$ 5.8
	Linear Regression (attribute 11 (PTRATIO))	61.03 $\pm$ 3.8	62.43 $\pm$ 7.8
	Linear Regression (attribute 12 (BLACK))	67.79 $\pm$ 4.2	71.18 $\pm$ 8.6
	Linear Regression (attribute 13 (LSTAT))	27.69 $\pm$ 1.8	26.45 $\pm$ 4.8
	Linear Regression (all attributes)	0.0074 $\pm$ 0.041	3.41 * 10 <sup>-5</sup> $\pm$ 6.1 * 10 <sup>-4</sup>
	Kernel Ridge Regression (all attributes)	0.012 $\pm$ 0.34	1.21 * 10 <sup>-5</sup> $\pm$ 4.3 * 10 <sup>-4</sup>

6. (a)

$$\mathcal{E}(f) = E[L_c(y, \hat{y})] = E[[y \neq \hat{y}]c_y] = \sum_{x \in X} \sum_{y \in Y} [y \neq f(x)]c_y p(x, y)$$

Applying Bayes rule, we can write:

$$\mathcal{E}(f) = \sum_{x \in X} \left[ \sum_{y \in Y} [y \neq f(x)]c_y p(y|x) \right] p(x)$$

Now let's find the value of the Bayes Estimator at a specific point  $x = x'$ :

$$\mathcal{E}(f(x')) = \left[ \sum_{y \in Y} [y \neq f(x')]c_y p(y|x') \right] p(x')$$

As we can see if for any  $y$ ,  $y \neq f(x)$  the cost would be  $c_y$ , hence we want to choose  $f(x')$  in a way that would increase the number of times we get  $y = f(x')$  to have more 0 costs, which would give us the minimum error. As a result, the Bayes Estimator  $f(x')$  is mode of the probability distribution.

(b)

$$\mathcal{E}(f) = E[L(y, \hat{y})] = E[|y - \hat{y}|] = \int |y - f(x)| dP(x, y)$$

In order to derive Bayes Estimator, we need to minimize the expected error. First we need to apply Bayes rule:

$$\mathcal{E}(f) = \int_{x \in X} \left\{ \int_{y \in Y} |y - f(x)| dP(y|x) \right\} dP(x)$$



Then we need to find the value of  $f(x')$  at a fixed point  $x=x'$ :

$$\begin{aligned}\mathcal{E}(f(x')) &= \int_{y \in Y} \{|y - f(x')| dP(y|x')\} dP(x') \\ &= \int_{y < f(x')} (f(x') - y) dP(y|x') dP(x') + \int_{y \geq f(x')} (y - f(x')) dP(y|x') dP(x')\end{aligned}$$

Now let's assume  $z = f(x')$ , we need to differentiate  $\mathcal{E}(f(x'))$  w.r.t  $z$  in order to find the Bayes Estimator:

$$\begin{aligned}\frac{\partial e}{\partial z} &= \int_{y < z} dP(y|x') - \int_{y \geq z} dP(y|x') = 0 \\ \int_{y < z} dP(y|x') &= \int_{y \geq z} dP(y|x')\end{aligned}$$

As  $P(y|x')$  is a distribution,

$$\int_{y \in Y} dP(y|x') = 1$$

Hence we can conclude that:

$$\int_{y < z} dP(y|x') = \int_{y \geq z} dP(y|x') = 1/2$$

So the Bayes Estimator is the median of the probability distribution  $P(y|x')$ .

7. (a) We have:

$$K_c(\mathbf{x}, \mathbf{z}) = c + \sum_{i=1}^n x_i z_i$$

where  $x, z \in \mathbb{R}^n$

$K_c$  would be positive semidefinite if and only if:

$$K_c(\mathbf{x}, \mathbf{z}) = \langle \phi(\mathbf{x}), \phi(\mathbf{z}) \rangle$$

if  $\phi(x) = (x_{i_1} x_{i_2} x_{i_3} \dots x_{i_n})$  then  $\sum_{i=1}^n x_i z_i = (x^T z)$  which is already a positive semidefinite by definition. So for any  $c \geq 0$ ,  $K_c$  is positive semidefinite.

(b)  $c$  will act as a bias term if  $K_c$  is used as kernel for linear regression.

$$\epsilon(f) = \sum_{i=1}^l (\sum_{j=1}^l \alpha_j (x_i, x_j) + c \alpha_j - y_i)^2 \text{ where } f \text{ is predictor.}$$

As you can see, in the expected error,  $c$  acts as a bias term when we use  $K_c$  for linear regression.

8. As  $\beta \rightarrow \infty$ , Gaussian kernel  $K_\beta$  starts to simulate 1-Nearest Neighbour. We illustrate this below by defining the classifier as follows,

$$f(x) = \begin{cases} -1 & \sum_{i=1}^m \alpha_i K_\beta(x_i, x) < 0 \\ 1 & \text{otherwise} \end{cases}$$

Since  $K_\beta(x_i, x_j) = \exp(-\beta \|x_i - x_j\|^2)$ , as  $\beta \rightarrow \infty$ ,  $K_\beta = 0$  if  $i \neq j$  and  $K_\beta = 1$  if  $i = j$ . Which means, if there are no duplicate points, all the other points except  $i = j$  becomes negligible.

From the notes,

$$\alpha_i = \frac{1}{2\lambda} V' \left( y_i, \sum_{j=1}^m \alpha_j K_\beta(\mathbf{x}_i, \mathbf{x}_j) \right)$$

so  $\alpha_i$  becomes  $\alpha_i = \frac{1}{2\lambda} V'(y_i, \alpha_i)$

$$\alpha_i = \frac{1}{\lambda} (\alpha_i - y_i)$$

hence  $\alpha_i = ky_i$

We also know that  $\forall i, j, \alpha_i = \alpha_j$

Now we can modify our classifier to be,

$$f(x) = \begin{cases} -1 & \sum_{i=1}^m ky_i K_\beta(x_i, x) < 0 \\ 1 & \text{otherwise} \end{cases}$$

Our classifier needs to predict  $y_p$  if  $x$  is closer to  $x_p$ . Which means,

$$|y_p K_\beta(\mathbf{x}_p, \mathbf{x})| > \sum_{i \in \{1, \dots, m\} \setminus \{p\}} |y_i K_\beta(\mathbf{x}_i, \mathbf{x})|$$

Since as  $\beta \rightarrow \infty$ ,  $|y_p K_\beta(\mathbf{x}_p, \mathbf{x}_p)| \rightarrow \infty$  and  $|y_i K_\beta(\mathbf{x}_i, \mathbf{x}_j)| \rightarrow 0$ , our classifier  $f(x)$  will predict  $y_p$  when  $x$  is closer to  $x_p$ . Hence as  $\beta \rightarrow \infty$ , our Gaussian kernel  $K_\beta$  simulate a 1-NN.

9. Let's represent the initial configuration of a Whack-A-Mole as a vector of 1s and 0s, which 1 represents a mole being present in the square and 0 represents an empty square without a mole. For example, for the configuration given in the question the initial representation of the grid would look like: INITIAL = (0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1)

As we want to empty the grid from any moles, we want our final configuration to be: FINAL = (0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)

We will also represent the moves that need to be made to solve the Whack-A-Mole problem as a vector of 1s and 0s, where 1s would represent whacking a mole and 0 not whacking a mole. This way, we can represent all the possible moves and its affects in a matrix where each row (i) would represent the affect of a hit in the ith square in the

grid and the columns of the matrix would represent the squares that are affected by that move.

For example:  $moves = \begin{bmatrix} 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}$

Now we can solve Whack-A-Mole by solving the below equation:

$$MOVES * X + INITIAL = FINAL \quad (1)$$

Where X represents the moves that we need to actually make to reach to the final configuration of all squares being 0. Which means that if there's a solution to this equation, there exist a series of moves that would give us the final configuration.

$$X = MOVES^{-1} * (FINAL - INITIAL) \quad (2)$$

Computational complexity of this equation is polynomial. Computing the inverse of a  $n * n$  matrix is  $O(n^3)$  and matrix multiplication is  $O(n^2)$ . So the overall complexity would be:  $O(n^3 + n^2) \approx O(n^3)$